

LOCAL SEARCH ALGORITHMS

Chapter 4

Outline



- Local search algorithms
 - Hill-climbing search
 - Simulated annealing search
 - Local beam search
 - Genetic algorithms

Local search algorithms

- In many optimization problems
 - ▣ the path from the start node to the goal is irrelevant
 - ▣ the goal state itself is the solution
- **State space** = set of "complete" configurations
- Find **configuration** satisfying constraints
- In such cases, we can use local search algorithms
 - ▣ keep a single "current" state
 - ▣ try to improve it

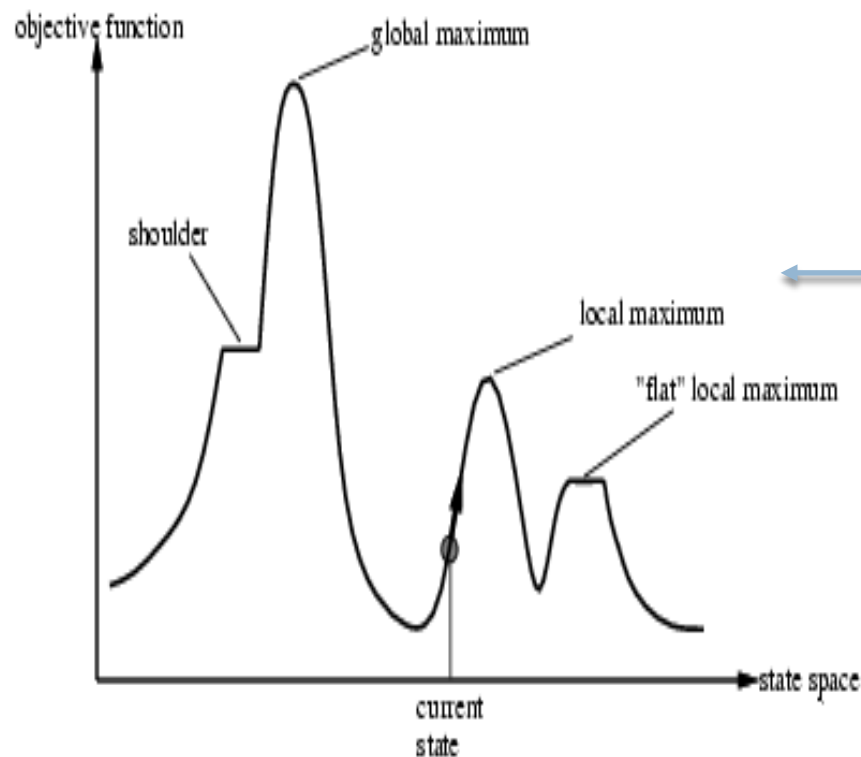
Example: n -queens

- Put n queens on an $n \times n$ board with no two queens on the **same** row, column, or diagonal
 - Each state has n queens on the board, one per column
 - Successors of a state: all possible states generated by moving a single queen to another square in the same column



State-space landscape

Local search algorithms explore the **state-space landscape**




A **landscape** has

- **location** (defined by the **state**)
- **elevation** (defined by the **value** of heuristic cost function or objective function)

The **aim** is to find:

- **a global minimum** (lowest valley)
if **elevation** corresponds to **cost**
- **a global maximum** (highest peak)
if **elevation** corresponds to **objective function**

Hill-climbing search



- Assume the **elevation** corresponds to the **objective function**
- **Hill-climbing search** modifies the **current state** to try to **improve it**

function **HILL-CLIMBING**(**problem**) **returns** a **state** that is a **local maximum**

current \leftarrow MAKE-NODE(problem.INITIAL-STATE)

loop do


neighbor \leftarrow a **highest-valued** successor of **current**

if **neighbor**.VALUE \leq **current**.VALUE **then return** **current**.STATE

current \leftarrow **neighbor**

- Picks a **neighbor** with the highest value
- Usually **chooses at random** among neighbors with maximum value
- Terminates when it reaches a "peak" where no neighbor has a higher value

Hill-climbing search



□ Hill climbing **often gets stuck** for the following reasons:


□ Local maxima: a **peak** that is

- **higher** than each of its neighboring states
- **lower** than the global maximum

□ Plateaux: a **flat area** of the state-space landscape

- **flat local maximum**, from which no progress is possible
- **shoulder**, from which progress is possible

Hill-climbing search



Eg.: 8-queens

Starting from a randomly generated 8-queens state

- **Hill climbing gets stuck 86%** of the time
- **Solving only 14%** of problem instances
- **It works quickly**, taking just **4 steps** on average when it succeeds and **3 steps** when it gets stuck, even if the **state space** with \approx **17 million states**

Allowing up to 100 consecutive sideways moves:

- **Solving 94%** of problem instances
- The algorithm averages **21 steps** for each successful instance and **64 steps** for each failure

Hill climbing variants

□ Stochastic hill climbing

- chooses at random from the set of all improving neighbors

□ First-choice hill climbing

- jumps to the first improving neighbor found

□ Random-restart hill climbing

- **series of hill climbing** runs until a goal is found
- It will find a good solution very quickly

Eg., For three million queens, it can find solutions in under a minute

Hill climbing and Random Walk

- **Hill-climbing algorithm** that never makes “downhill” moves toward states with lower value is **incomplete**
 - because it can get stuck on a local maximum
- **Purely random walk**: moving to a **successor** chosen **uniformly at random** from the set of successors is **complete** but extremely **inefficient**
- **Idea**: to **combine** hill climbing with a **random walk** that yields both **efficiency** and **completeness**

Simulated annealing search

- **Simulated annealing:** a version of stochastic hill climbing where some **downhill moves** are allowed
 - If the move improves the situation → **always accepted**
 - Otherwise → **accepted** with **some probability** less than 1

- Downhill moves
 - **accepted early** in the annealing schedule
 - then **accepted less** often **as time goes on**

Simulated annealing search

Idea: **escape local maxima** by allowing some "bad" moves but gradually decrease their frequency

function **SIMULATED-ANNEALING**(problem, schedule) **returns** a solution state

inputs: problem, a problem

schedule, a mapping from time to "temperature"

current \leftarrow MAKE-NODE(problem.INITIAL-STATE)

for $t = 1$ **to** ∞ **do**

$T \leftarrow$ schedule(t)

if $T = 0$ **then return** current

next \leftarrow a randomly selected successor of current

$\Delta E \leftarrow$ next.VALUE $-$ current.VALUE

if $\Delta E > 0$ **then** current \leftarrow next

else current \leftarrow next only with probability $e^{(\Delta E/T)}$

The **schedule input** determines the value of the temperature T as a function of time

Simulated annealing search

Idea: **escape local maxima** by allowing some "bad" moves but gradually decrease their frequency

function **SIMULATED-ANNEALING**(problem, schedule) **returns** a solution state

inputs: problem, a problem

schedule, a mapping from time to "temperature"

current \leftarrow MAKE-NODE(problem.INITIAL-STATE)

for $t = 1$ **to** ∞ **do**

$T \leftarrow$ schedule(t)

if $T = 0$ **then return** current

next \leftarrow a randomly selected successor of current

$\Delta E \leftarrow$ next.VALUE $-$ current.VALUE

if $\Delta E > 0$ **then** current \leftarrow next

else current \leftarrow next only with probability $e^{(\Delta E/T)}$

The higher the temperature
the higher the probability of
making a non-improving move

Properties of simulated annealing search

- One can prove: If T decreases slowly enough,
then **simulated annealing search** will find a global optimum
with probability approaching 1
- Simulated annealing widely used in
 - ▣ airline scheduling
 - ▣ large-scale optimization tasks etc

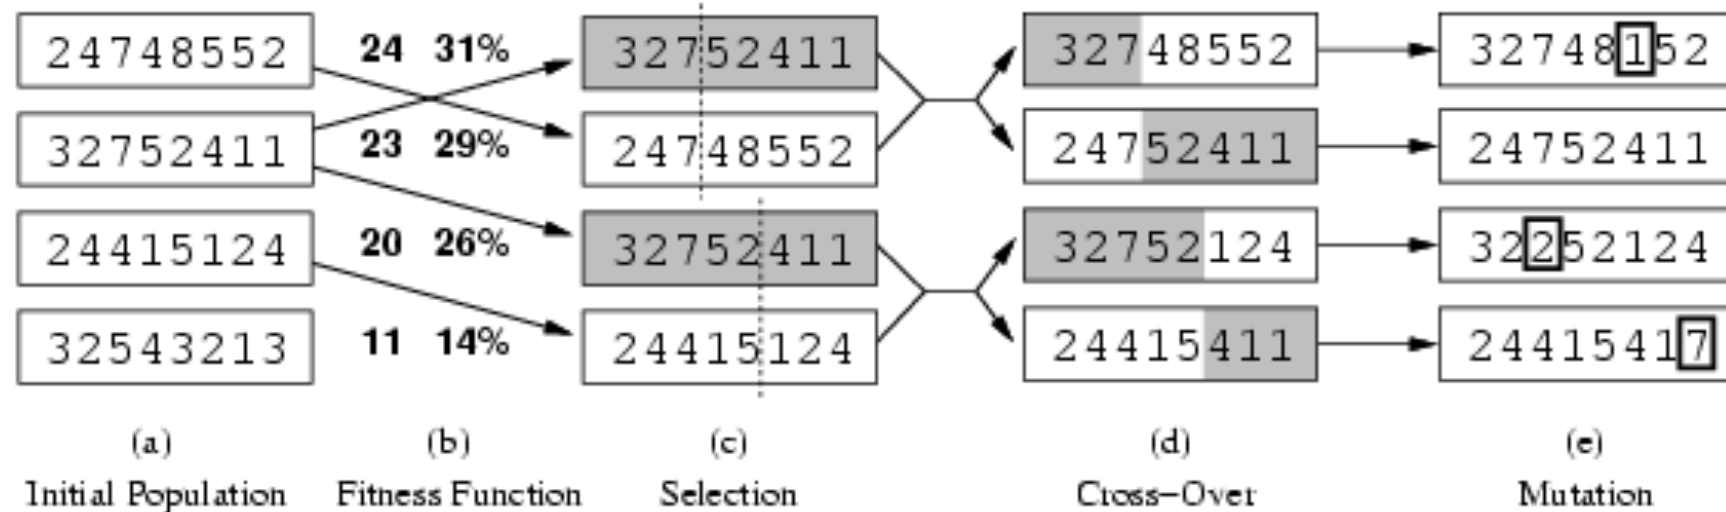

Local beam search

- **Keep track** of **k states** rather than just one
- **Start** with **k randomly** generated **states**
- **At each iteration** all **the successors of** all **k states** are generated
 - **If** any one is a **goal state**, algorithm **stops**
 - **Else select** the **k best successors** from the complete list and **repeat** (they could be all successors of the same node)

Genetic algorithms

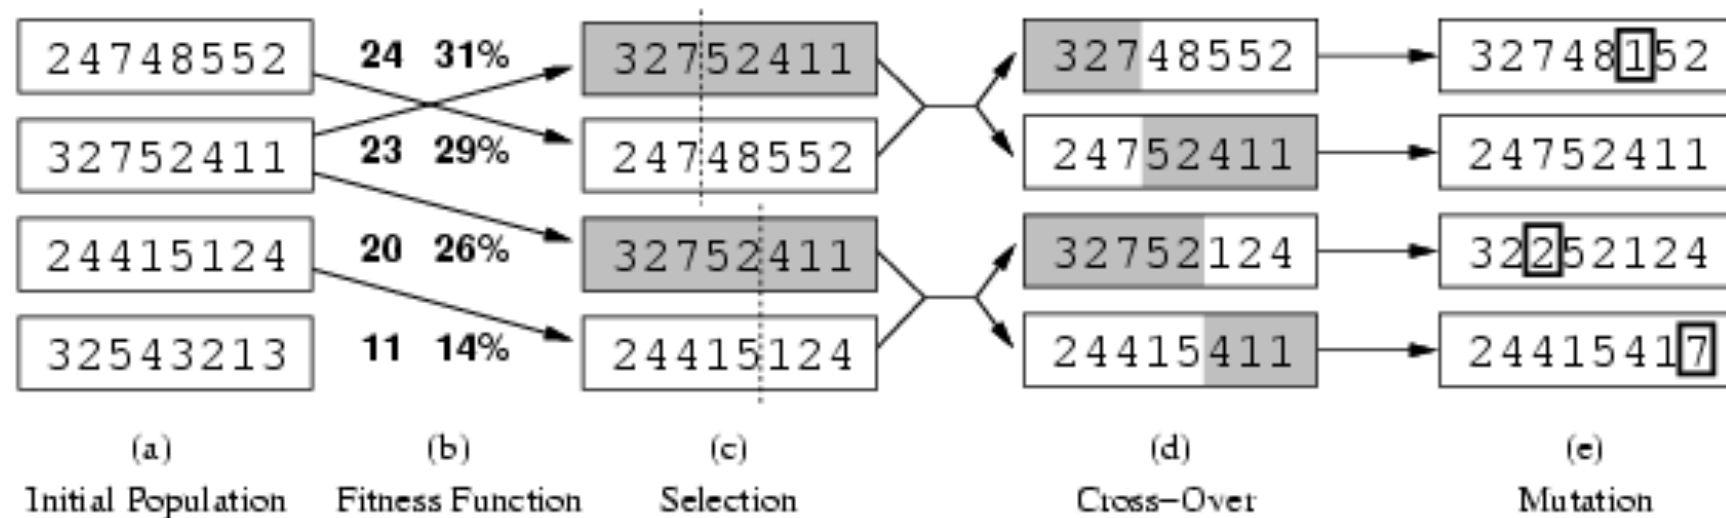

- A **successor state** is generated by combining two parent states rather than by modifying a single state
- **Start** with **k randomly** generated **states** (**population**)
 - A **state** is represented **as a string** over a finite alphabet (often a string of 0s and 1s or digits)
- **Each state** is associated to **a value** via an **evaluation function** (**fitness function**)
 - **Returns higher** values for **better** states
- **The next generation of states** is produced by selection, crossover, and mutation

Genetic algorithms




- States: strings of 8 digits representing **8-queens states**
- Fitness function:
 - High values for better states
 - We use the **number of non-attacking pairs of queens**
 - The higher the fitness the more likely the node is to be **selected** for reproductions

Genetic algorithms



- For each pair to be mated, a **crossover** point is chosen randomly from the positions in the string
- Each location is subject to random **mutation** with a small independent probability



The 8-queens **states** corresponding to:
the **first two parents** in (c)

the **first child** in (d)

